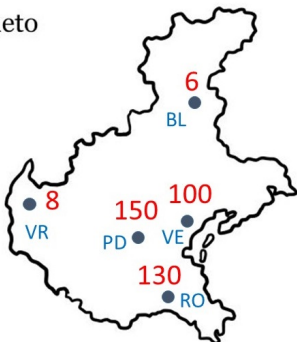


HOMEWORK 2

k-center with outliers: motivation

Veneto



Red numbers = potential customers

If you must open 3 stores, where would you open them?

If you used k-center (with $k=3$) to select the store locations, what would be the result?

k-center with outliers: problem

Input Set P of N points from (M, d) , integers $k, z \in [1, N]$

Output Set $S \subset P$ of k centers which minimize

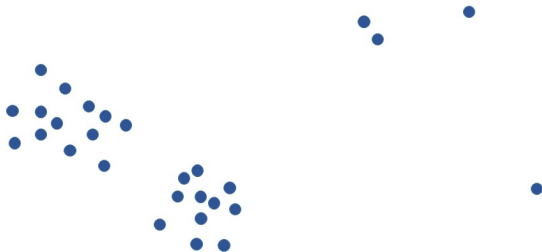
$$\Phi_{\text{kcenter}}(P - Z_S, S) = \max_{x \in P - Z_S} d(x, S)$$

where $Z_S = z$ points of P farthest from S .

Remarks:


- It is the k-center problem where the objective function is allowed to disregard the distances of the z points farthest from the center. For noisy datasets, this yields a better center selection.
- The objective function is uniquely determined by P and S
- with $z=0$ we have the standard k-center problem

Example



Input P

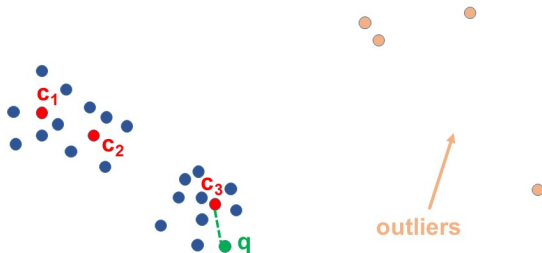
Example



Optimal solution with $k=3$ centers

Value of objective function = $d(c_1, q)$

Example



Optimal solution with $k=3$ centers and $z=4$ outliers

Value of objective function = $d(c_3, q)$

typically much less than the value of the objective function with $z=0$

k-center with outliers: weighted variant

Input Set $(P, w(\cdot))$ of N weighted points from (M, d) , integers $k, z \in [1, N]$

Output Set $S \subset P$ of k centers which minimize

$$\Phi_{\text{kcenter}}(P - Z_S, S) = \max_{x \in P - Z_S} d(x, S)$$

where

$Z_S =$ largest set of points farthest from S of total weight $\leq z$

Remark: if we consider P sorted by distance from S , then Z_S is the longest suffix of points, in the sorted sequence, of total weight $\leq z$.

Example

	Points sorted by distance from S											
distance	0	0	0	0	1.3	1.6	1.9	2.4	2.5	2.7	4.1	6.3
weight	9	11	6	4	13	2	12	3	4	2	1	2

$z = 3, 4$
value of obj. function with $z = 3, 4$

$z = 9, 10, 11$
value of obj. function with $z = 9, 10, 11$

k-center with outliers: algorithm kcenterOUT

For a radius $r > 0$, a set $Z \subset P$, and a point x , define the ball of Z with radius r centered at x as

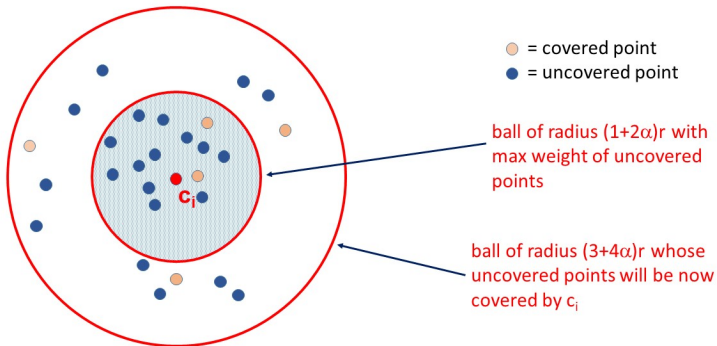
$$B_Z(x, r) = \{y \in Z : d(x, y) \leq r\}.$$

Overview of the algorithm:

- Geometric sequence of guesses $r = r_{\min} \cdot 2^j$ for $j = 0, 1, \dots$, where r_{\min} is a lower bound to the optimal objective function.
- S = current centers, Z = uncovered points, α = parameter.
- For each guess r , initialize $S = \emptyset$ and $Z = P$, and execute k iterations. In the i -th iteration:
 - add to S the point $c_i \in P$ which maximizes the total weight in $B_Z(c_i, (1 + 2\alpha)r)$.
 - remove from Z the points of $B_Z(c_i, (3 + 4\alpha)r)$
- The algorithm stops at the smallest guess r for which the points of Z (i.e., the outliers), have total weight $\leq z$, and returns S

when the total weight of the points of Z is $> z \Rightarrow$
proceed with the next guess $2 \cdot r$

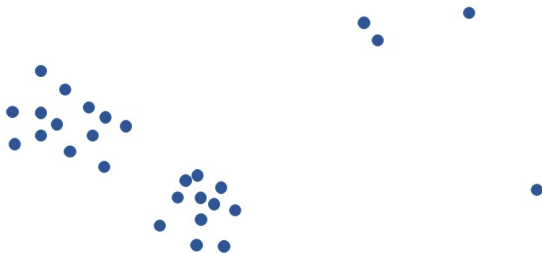
Selection of i -th center



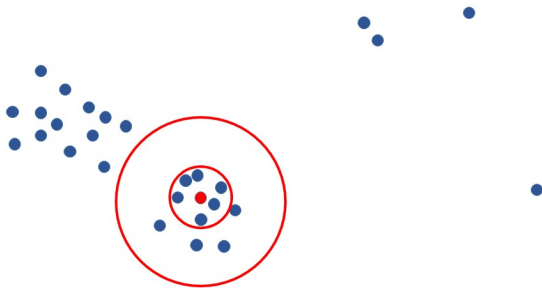
kcenterOUT(P, k, z, α): pseudocode

```
 $r \leftarrow (\text{min distance between first } k + z + 1 \text{ points})/2; \text{ // } \tau_{\min}$   
while (true) do  
   $Z \leftarrow P; S \leftarrow \emptyset; W_Z = \sum_{x \in P} w(x);$   
  while (( $|S| < k$ ) AND ( $W_Z > 0$ )) do  
     $\text{max} \leftarrow 0;$   
    foreach  $x \in P$  do  
       $\text{ball-weight} \leftarrow \sum_{y \in B_Z(x, (1+2\alpha)r)} w(y);$   
      if ( $\text{ball-weight} > \text{max}$ ) then  
         $\text{max} \leftarrow \text{ball-weight};$   
         $\text{newcenter} \leftarrow x;$   
     $S \leftarrow S \cup \{\text{newcenter}\};$   
    foreach ( $y \in B_Z(\text{newcenter}, (3 + 4\alpha)r)$ ) do  
      remove  $y$  from  $Z;$   
      subtract  $w(y)$  from  $W_Z;$   
  if ( $W_Z \leq z$ ) then return  $S;$   
  else  $r \leftarrow 2r;$ 
```


Example: $k = 2$, $z = 4$, unit weights



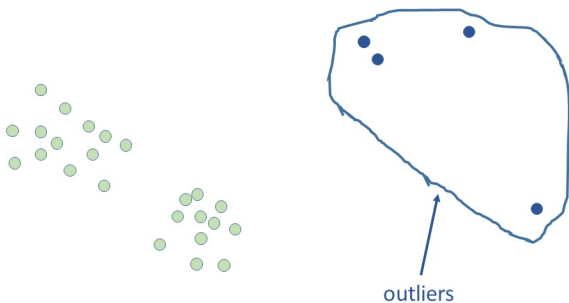
Example: $k = 2$, $z = 4$, unit weights



Example: $k = 2$, $z = 4$, unit weights



Example: $k = 2$, $z = 4$, unit weights



Quality of solution

- The value of the objective function for the returned solution S is at most $(3 + 4\alpha)r$, where r is the last guess. However, it can be smaller.
- If `kcenterOUT` is executed on the entire dataset with unit weights and $\alpha = 0$, then it returns a 3-approximation.
- If `kcenterOUT` is executed on a weighted coreset $T \subseteq P$ obtained by using $k + z$ -center on P (or on each of L partitions of P), with $\alpha = 2$, then it returns a 13-approximation.

Homework 2

Write a sequential program that:

- 1 Contains a method that implements sequentially `kcenterOUT` (**weighted variant and Euclidean points!**)
- 2 Reads P, k, z input parameters and runs `kcenterOUT($P, k, z, \alpha = 0$)` using unit weights.
- 3 Computes the value of the objective function.
- 4 Reports several statistics.

Detailed specification in Moodle Exam

Homework 2: hints

- Representation of points (arbitrary dimensionality): *Each point is an instance of*
 - **Java:** class `org.apache.spark.mllib.linalg.Vector`, which can be manipulated through static methods offered by class `org.apache.spark.mllib.linalg.Vectors`.
 - **Python:** tuple of float
- You will not use Spark (except for the points representation in Java)
- Distance computations are the bottleneck (especially in Python!). You may precompute all distances: up to 10^4 they are likely to fit it RAM, and for the MapReduce implementation (Homework 3) we won't need to run it on larger instances.
- Carefully select the representation for Z .
- Debug your code for correctness and efficiency (try to identify sources of inefficiency!)

Homework 3

In Homework 3 you will implement a 2-round MR-algorithm for k -center with outliers, where Round 1 computes a **weighted coreset** of size $(k + z) \cdot L$ using Farthest-First Traversal in each of L partitions, and Round 2 runs `kcenterOUT` on the weighted coreset.

We will provide a template for the homework and you must:

- 1 write the code for FFT (Round 1);
- 2 add code for `kcenterOUT` (Round 2), recycled from Homework 2;
- 3 run the algorithm on CloudVeneto, assessing running time, scalability with respect to parallelism, and quality of solution.

Details will be given soon.